

Recurrent Neural Networks (RNNs)

RNNs are a class of neural networks designed to process sequential data. Unlike feedforward networks,

RNNs maintain a hidden state that allows them to capture temporal dependencies.

They are widely used in tasks such as language modeling, speech recognition, and time-series forecasting. However, RNNs suffer from vanishing gradients, which led to the development of LSTMs and GRUs for improved performance.